

Macro-F1 across zero-shot baselines and fine-tuned models
(5-fold CV mean $\pm$ std where applicable)

